

More on grouping and aggregation

You've discovered that pandas is a Python library that facilitates reviewing and manipulating tabular data. In addition, `groupby()` and `agg()` are essential `DataFrame` methods that data professionals use to group, aggregate, summarize, and better understand data. In this reading, you'll review how these functions work, as well as when and how to apply them.

`groupby()`

The `groupby()` function is a method that belongs to the `DataFrame` class. It works by splitting data into groups based on specified criteria, applying a function to each group independently, then combining the results into a data structure. When applied to a dataframe, the function returns a groupby object. This groupby object serves as the foundation for different data manipulation operations, including:

- Aggregation: Computing summary statistics for each group
- Transformation: Applying functions to each group and returning modified data
- Filtration: Selecting specific groups based on certain conditions
- Iteration: Iterating over groups or values

Here are some examples that use the `groupby()` function on a dataframe consisting of different articles of clothing:

```
1
2
3
4
5
6
7
clothes = pd.DataFrame({'type': ['pants', 'shirt', 'shirt', 'pants', 'shirt', 'pants'],
                          'color': ['red', 'blue', 'green', 'blue', 'green', 'red'],
                          'price_usd': [20, 35, 50, 40, 100, 75],
                          'mass_g': [125, 440, 680, 200, 395, 485]})
```

`clothes`

RunReset

	color	mass_g	price_usd	type
0	red	125	20	pants
1	blue	440	35	shirt
2	green	680	50	shirt
3	blue	200	40	pants
4	green	395	100	shirt
5	red	485	75	pants

Grouping the dataframe by `type` results in a `DataFrameGroupBy` object:

1
2

3

```
grouped = clothes.groupby('type')
print(grouped)
print(type(grouped))
```

RunReset

```
<pandas.core.groupby.DataFrameGroupBy object at 0x7fc9d6323160>
<class 'pandas.core.groupby.DataFrameGroupBy'>
```

However, an aggregation function can be applied to the groupby object:

1

2

```
grouped = clothes.groupby('type')
grouped.mean()
```

RunReset

```
      mass_g  price_usd
type
pants    270.0   45.000000
shirt    505.0   61.666667
```

In the preceding example, `groupby()` combined all the items into groups based on their type and returned a `DataFrame` object containing the mean of each group for each numeric column in the dataframe. Note: In future versions of pandas it will be necessary to specify a `numeric_only` parameter when applying certain aggregation functions—like `mean`—to a groupby object. `numeric_only` refers to the datatype of each column. In earlier versions of pandas (like the version on this platform) it isn't necessary to specify `numeric_only=True`, but in future versions this must be done. Otherwise, it will be necessary to indicate the specific columns to be captured.)

In addition, groups may be created based on multiple columns:

1

2

```
clothes.groupby(['type', 'color']).min()
```

RunReset

```
      mass_g  price_usd
type  color
pants blue      200        40
      red      125        20
shirt blue      440        35
      green     395        50
```

In the preceding example, `groupby()` was called directly on the clothes dataframe. The data was grouped first by `type`, then by `color`. This resulted in four groups—the number of different existing combinations of values for type and color. Then, the `min()` function was applied to the result to filter each group by its minimum value.

To simply return the number of observations there are in each group, use the `size()` method. This will result in a `series` object with the relevant information:

1

```
clothes.groupby(['type', 'color']).size()
```

```
RunReset
```

```
type    color
pants  blue      1
       red       2
shirt  blue      1
       green     2
dtype: int64
```

Built-in aggregation functions

The previous examples demonstrated the `mean()`, `min()`, and `size()` aggregation functions applied to groupby objects. There are many available built-in aggregation functions. Some of the more commonly used include:

- `count()`: The number of non-null values in each group
- `sum()`: The sum of values in each group
- `mean()`: The mean of values in each group
- `median()`: The median of values in each group
- `min()`: The minimum value in each group
- `max()`: The maximum value in each group
- `std()`: The standard deviation of values in each group
- `var()`: The variance of values in each group

agg()

The [`agg\(\)`](#) function is useful when you want to apply multiple functions to a dataframe at the same time. `agg()` is a method that belongs to the `DataFrame` class. It stands for “aggregate.” Its most important parameters are:

- `func`: The function to be applied
- `axis`: The axis over which to apply the function (default= 0).

Following are some examples of how `agg()` can be used. Note that they demonstrate how this function can be used by itself (without `groupby()`). Note also that, due to platform limitations, some of the following code blocks are not executable. In these cases, output is provided as an image. Here is the original `clothes` dataframe again as a reminder:

```
clothes
```

```
RunReset
```

```
   color  mass_g  price_usd  type
0   red     125         20  pants
1  blue     440         35  shirt
2 green     680         50  shirt
3  blue     200         40  pants
4 green     395        100  shirt
5   red     485         75  pants
```

The following example applies the `sum()` and `mean()` functions to the `price` and `mass_g` columns of the `clothes` dataframe.

1

```
clothes[['price_usd', 'mass_g']].agg(['sum', 'mean'])
```

Output:

	price_usd	mass_g
sum	320.000000	2325.0
mean	53.333333	387.5

Notice the following:

- The two columns are subset from the dataframe before applying the `agg()` method. If you don't subset the relevant columns first, `agg()` will attempt to apply `sum()` and `mean()` to *all* of the columns, which wouldn't work because some columns contain strings. (Technically, `sum()` would work, but it would return something useless because it would just combine all the strings into one long string.)
- The `sum()` and `mean()` functions are entered as strings in a list, without their parentheses. This will work for any built-in aggregation function.

In this next example, different functions are applied to different columns.

3

1

2

```
    })
clothes.agg({'price_usd': 'sum',
            'mass_g': ['mean', 'median']})
```

Output:

	price_usd	mass_g
sum	320.0	NaN
mean	NaN	387.5
median	NaN	417.5

Notice the following:

- Columns are not subset from the dataframe before applying the `agg()` function. This is unnecessary because the columns are specified within the `agg()` function itself.
- The argument to the `agg()` function is a dictionary whose keys are columns and whose values are the functions to be applied to those columns. If multiple functions are applied to a column, they are entered as a list. Again, each built-in function is entered as a string without parentheses.

- The resulting dataframe contains **NaN** values where a given function was not designated to be used.

The following example applies the `sum()` and `mean()` functions across axis 1. In other words, instead of applying the functions down each column, they're applied over each row.

1

```
clothes[['price_usd', 'mass_g']].agg(['sum', 'mean'], axis=1)
```

Output:

	sum	mean
0	145.0	72.5
1	475.0	237.5
2	730.0	365.0
3	240.0	120.0
4	495.0	247.5
5	560.0	280.0

groupby() with agg()

The `groupby()` and `agg()` functions are often used together. In such cases, first apply the `groupby()` function to a dataframe, then apply the `agg()` function to the result of the `groupby()`. For reference, here is the `clothes` dataframe once again.

1

2

clothes

RunReset

	color	mass_g	price_usd	type
0	red	125	20	pants
1	blue	440	35	shirt
2	green	680	50	shirt
3	blue	200	40	pants
4	green	395	100	shirt
5	red	485	75	pants

In the following example, the items in `clothes` are grouped by `color`, then each of those groups has the `mean()` and `max()` functions applied to them at the `price_usd` and `mass_g` columns.

1

2

3

```
clothes.groupby('color').agg({'price_usd': ['mean', 'max'],
                             'mass_g': ['mean', 'max']})
```

RunReset

	price_usd		mass_g	
	mean	max	mean	max
color				
blue	37.5	40	320.0	440
green	75.0	100	537.5	680
red	47.5	75	305.0	485

MultiIndex

You might have noticed that, when functions are applied to a groupby object, the resulting dataframe has tiered indices. This is an example of **MultiIndex**. MultiIndex is a hierarchical system of dataframe indexing. It enables you to store and manipulate data with any number of dimensions in lower dimensional data structures such as series and dataframes. This facilitates complex data manipulation.

This course will not require any deep knowledge of hierarchical indexing, but it's helpful to be familiar with it. Consider the following example:

1
2

```
grouped = clothes.groupby(['color', 'type']).agg(['mean', 'min'])
grouped
```

RunReset

		mass_g		price_usd	
		mean	min	mean	min
color	type				
blue	pants	200.0	200	40.0	40
	shirt	440.0	440	35.0	35
green	shirt	537.5	395	75.0	50
red	pants	305.0	125	47.5	20

Notice that **color** and **type** are positioned lower than the column names in the output. This indicates that **color** and **type** are no longer columns, but named row indices. Similarly, notice that **price_usd** and **mass_g** are positioned above **mean** and **min** in the output of column names, indicating a hierarchical column index.

If you inspect the row index, you'll get a **MultiIndex** object containing information about the row indices:

1
2

```
grouped.index
```

RunReset

```
MultiIndex(levels=[['blue', 'green', 'red'], ['pants', 'shirt']],
            labels=[[0, 0, 1, 2], [0, 1, 1, 0]],
            names=['color', 'type'])
```

The column index shows a **MultiIndex** object containing information about the column indices:

1
2

`grouped.columns`

RunReset

```
MultiIndex(levels=[['mass_g', 'price_usd'], ['mean', 'min']],
            labels=[[0, 0, 1, 1], [0, 1, 0, 1]])
```

To perform selection on a dataframe with a MultiIndex, use `loc[]` selection and put indices in parentheses. Here are some examples on `grouped`, which is a dataframe with a two-level row index and a two-level column index. For reference, here is the `grouped` dataframe:

1
2

`grouped`

RunReset

		mass_g		price_usd	
		mean	min	mean	min
color	type				
blue	pants	200.0	200	40.0	40
	shirt	440.0	440	35.0	35
green	shirt	537.5	395	75.0	50
red	pants	305.0	125	47.5	20

To select a first-level (top) column:

1
2

`grouped.loc[:, 'price_usd']`

RunReset

		mean	min
color	type		
blue	pants	40.0	40
	shirt	35.0	35
green	shirt	75.0	50
red	pants	47.5	20

To select a second-level (bottom) column:

1
2

`grouped.loc[:, ('price_usd', 'min')]`

RunReset

color	type	
blue	pants	40
	shirt	35
green	shirt	50
red	pants	20

Name: (price_usd, min), dtype: int64

To select first-level (left-most) row:

1
2

```
grouped.loc['blue', :]
```

RunReset

	mass_g		price_usd	
type	mean	min	mean	min
pants	200.0	200	40.0	40
shirt	440.0	440	35.0	35

To select a bottom-level (right-most) row:

1
2

```
grouped.loc[('green', 'shirt'), :]
```

RunReset

mass_g	mean	537.5
	min	395.0
price_usd	mean	75.0
	min	50.0

Name: (green, shirt), dtype: float64

And you can even select individual values:

1
2

```
grouped.loc[('blue', 'shirt'), ('mass_g', 'mean')]
```

RunReset

440.0

If you want to remove the row MultiIndex from a groupby result, include `as_index=False` as a parameter to your `groupby()` statement:

1
2

```
clothes.groupby(['color', 'type'], as_index=False).mean()
```

RunReset

	color	type	mass_g	price_usd
0	blue	pants	200.0	40.0
1	blue	shirt	440.0	35.0
2	green	shirt	537.5	75.0
3	red	pants	305.0	47.5

Notice how `color` and `type` are no longer row indices, but named columns. The row indices are the standard enumeration beginning from zero.

Again, you will not be expected to do any complex manipulations of hierarchically indexed data in this course, but it's helpful to have a basic understanding of how MultiIndex works, especially because `groupby()` manipulations typically result in a MultiIndex dataframe by default.

Key takeaways

`groupby()` will be an essential function in your work as a data professional, as it enables efficient combining and analysis of data. Similarly, `agg()` will help you apply multiple functions dynamically across a specified axis of a dataframe. Either on their own or when used together, these tools give data professionals deep access to data and help bring about successful projects.